


# End-to-end pipeline; the C1 rule replaces the Pareto search



Every transformer is fit on the training fold only and applied unchanged to the test fold. In the primary arm the whole front-end is refit at each  $N$ , so no representation ever sees more data than the classifier it feeds.